

DataFrames

DataFrames are the workhorse of pandas and are directly inspired by the R programming language. We can think of a DataFrame as a bunch of Series objects put together to share the same index. Let's use pandas to explore this topic!

```
In [1]: 1 import pandas as pd
        2 import numpy as np
```

```
In [2]: 1 from numpy.random import randn
        2 np.random.seed(101)
```

```
In [3]: 1 df = pd.DataFrame(randn(5,4),index='A B C D E'.split(),columns='W X Y
```

```
In [4]: 1 df
```

Out[4]:

	W	X	Y	Z
A	2.706850	0.628133	0.907969	0.503826
B	0.651118	-0.319318	-0.848077	0.605965
C	-2.018168	0.740122	0.528813	-0.589001
D	0.188695	-0.758872	-0.933237	0.955057
E	0.190794	1.978757	2.605967	0.683509

Selection and Indexing

Let's learn the various methods to grab data from a DataFrame

```
In [5]: 1 df['W']
```

Out[5]:

A	2.706850
B	0.651118
C	-2.018168
D	0.188695
E	0.190794

Name: W, dtype: float64

```
In [6]: 1 # Pass a list of column names
        2 df[['W', 'Z']]
```

Out[6]:

	W	Z
A	2.706850	0.503826
B	0.651118	0.605965
C	-2.018168	-0.589001
D	0.188695	0.955057
E	0.190794	0.683509

```
In [7]: 1 # SQL Syntax (NOT RECOMMENDED!)
        2 df.W
```

Out[7]: A 2.706850
B 0.651118
C -2.018168
D 0.188695
E 0.190794
Name: W, dtype: float64

DataFrame Columns are just Series

```
In [8]: 1 type(df['W'])
```

Out[8]: pandas.core.series.Series

Creating a new column:

```
In [9]: 1 df['new'] = df['W'] + df['Y']
```

```
In [10]: 1 df
```

Out[10]:

	W	X	Y	Z	new
A	2.706850	0.628133	0.907969	0.503826	3.614819
B	0.651118	-0.319318	-0.848077	0.605965	-0.196959
C	-2.018168	0.740122	0.528813	-0.589001	-1.489355
D	0.188695	-0.758872	-0.933237	0.955057	-0.744542
E	0.190794	1.978757	2.605967	0.683509	2.796762

Removing Columns

```
In [11]: 1 df.drop('new',axis=1)
```

Out[11]:

	W	X	Y	Z
A	2.706850	0.628133	0.907969	0.503826
B	0.651118	-0.319318	-0.848077	0.605965
C	-2.018168	0.740122	0.528813	-0.589001
D	0.188695	-0.758872	-0.933237	0.955057
E	0.190794	1.978757	2.605967	0.683509

```
In [12]: 1 # Not inplace unless specified!  
2 df
```

Out[12]:

	W	X	Y	Z	new
A	2.706850	0.628133	0.907969	0.503826	3.614819
B	0.651118	-0.319318	-0.848077	0.605965	-0.196959
C	-2.018168	0.740122	0.528813	-0.589001	-1.489355
D	0.188695	-0.758872	-0.933237	0.955057	-0.744542
E	0.190794	1.978757	2.605967	0.683509	2.796762

```
In [13]: 1 df.drop('new',axis=1,inplace=True)
```

```
In [14]: 1 df
```

Out[14]:

	W	X	Y	Z
A	2.706850	0.628133	0.907969	0.503826
B	0.651118	-0.319318	-0.848077	0.605965
C	-2.018168	0.740122	0.528813	-0.589001
D	0.188695	-0.758872	-0.933237	0.955057
E	0.190794	1.978757	2.605967	0.683509

Can also drop rows this way:

```
In [15]: 1 df.drop('E',axis=0)
```

Out[15]:

	W	X	Y	Z
A	2.706850	0.628133	0.907969	0.503826
B	0.651118	-0.319318	-0.848077	0.605965
C	-2.018168	0.740122	0.528813	-0.589001
D	0.188695	-0.758872	-0.933237	0.955057

Selecting Rows

```
In [16]: 1 df.loc['A']
```

Out[16]: W 2.706850
X 0.628133
Y 0.907969
Z 0.503826
Name: A, dtype: float64

Or select based on the position instead of label

```
In [17]: 1 df.iloc[2]
```

Out[17]: W -2.018168
X 0.740122
Y 0.528813
Z -0.589001
Name: C, dtype: float64

Selecting subset of rows and columns

```
In [18]: 1 df.loc['B','Y']
```

Out[18]: -0.8480769834036315

```
In [19]: 1 df.loc[['A','B'],['W','Y']]
```

Out[19]:

	W	Y
A	2.706850	0.907969
B	0.651118	-0.848077

Conditional Selection

An important feature of pandas is conditional selection using bracket notation, very similar to numpy:

```
In [20]: 1 df
```

Out[20]:

	W	X	Y	Z
A	2.706850	0.628133	0.907969	0.503826
B	0.651118	-0.319318	-0.848077	0.605965
C	-2.018168	0.740122	0.528813	-0.589001
D	0.188695	-0.758872	-0.933237	0.955057
E	0.190794	1.978757	2.605967	0.683509

```
In [21]: 1 df>0
```

Out[21]:

	W	X	Y	Z
A	True	True	True	True
B	True	False	False	True
C	False	True	True	False
D	True	False	False	True
E	True	True	True	True

```
In [22]: 1 df[df>0]
```

Out[22]:

	W	X	Y	Z
A	2.706850	0.628133	0.907969	0.503826
B	0.651118	NaN	NaN	0.605965
C	NaN	0.740122	0.528813	NaN
D	0.188695	NaN	NaN	0.955057
E	0.190794	1.978757	2.605967	0.683509

```
In [23]: 1 df[df['W']>0]
```

Out[23]:

	W	X	Y	Z
A	2.706850	0.628133	0.907969	0.503826
B	0.651118	-0.319318	-0.848077	0.605965
D	0.188695	-0.758872	-0.933237	0.955057
E	0.190794	1.978757	2.605967	0.683509

```
In [24]: 1 df[df['W']>0]['Y']
```

Out[24]: A 0.907969
B -0.848077
D -0.933237
E 2.605967
Name: Y, dtype: float64

```
In [25]: 1 df[df['W']>0][['Y', 'X']]
```

Out[25]:

	Y	X
A	0.907969	0.628133
B	-0.848077	-0.319318
D	-0.933237	-0.758872
E	2.605967	1.978757

For two conditions you can use | and & with parenthesis:

```
In [26]: 1 df[(df['W']>0) & (df['Y'] > 1)]
```

Out[26]:

	W	X	Y	Z
E	0.190794	1.978757	2.605967	0.683509

More Index Details

Let's discuss some more features of indexing, including resetting the index or setting it something else. We'll also talk about index hierarchy!

In [27]:

```
1 df
```

Out[27]:

	W	X	Y	Z
A	2.706850	0.628133	0.907969	0.503826
B	0.651118	-0.319318	-0.848077	0.605965
C	-2.018168	0.740122	0.528813	-0.589001
D	0.188695	-0.758872	-0.933237	0.955057
E	0.190794	1.978757	2.605967	0.683509

In [28]:

```
1 # Reset to default 0,1...n index
2 df.reset_index()
```

Out[28]:

	index	W	X	Y	Z
0	A	2.706850	0.628133	0.907969	0.503826
1	B	0.651118	-0.319318	-0.848077	0.605965
2	C	-2.018168	0.740122	0.528813	-0.589001
3	D	0.188695	-0.758872	-0.933237	0.955057
4	E	0.190794	1.978757	2.605967	0.683509

In [29]:

```
1 newind = 'CA NY WY OR CO'.split()
```

In [30]:

```
1 df['States'] = newind
```

In [31]:

```
1 df
```

Out[31]:

	W	X	Y	Z	States
A	2.706850	0.628133	0.907969	0.503826	CA
B	0.651118	-0.319318	-0.848077	0.605965	NY
C	-2.018168	0.740122	0.528813	-0.589001	WY
D	0.188695	-0.758872	-0.933237	0.955057	OR
E	0.190794	1.978757	2.605967	0.683509	CO

```
In [32]: 1 df.set_index('States')
```

Out[32]:

	W	X	Y	Z
States				
CA	2.706850	0.628133	0.907969	0.503826
NY	0.651118	-0.319318	-0.848077	0.605965
WY	-2.018168	0.740122	0.528813	-0.589001
OR	0.188695	-0.758872	-0.933237	0.955057
CO	0.190794	1.978757	2.605967	0.683509

```
In [33]: 1 df
```

Out[33]:

	W	X	Y	Z	States
A	2.706850	0.628133	0.907969	0.503826	CA
B	0.651118	-0.319318	-0.848077	0.605965	NY
C	-2.018168	0.740122	0.528813	-0.589001	WY
D	0.188695	-0.758872	-0.933237	0.955057	OR
E	0.190794	1.978757	2.605967	0.683509	CO

```
In [34]: 1 df.set_index('States',inplace=True)
```

```
In [35]: 1 df
```

Out[35]:

	W	X	Y	Z
States				
CA	2.706850	0.628133	0.907969	0.503826
NY	0.651118	-0.319318	-0.848077	0.605965
WY	-2.018168	0.740122	0.528813	-0.589001
OR	0.188695	-0.758872	-0.933237	0.955057
CO	0.190794	1.978757	2.605967	0.683509

Multi-Index and Index Hierarchy

Let us go over how to work with Multi-Index, first we'll create a quick example of what a Multi-Indexed DataFrame would look like:


```
In [36]: 1 # Index Levels
2 outside = ['G1', 'G1', 'G1', 'G2', 'G2', 'G2']
3 inside = [1,2,3,1,2,3]
4 hier_index = list(zip(outside,inside))
5 hier_index = pd.MultiIndex.from_tuples(hier_index)
```

```
In [37]: 1 hier_index
```

```
Out[37]: MultiIndex([( 'G1', 1),
                    ( 'G1', 2),
                    ( 'G1', 3),
                    ( 'G2', 1),
                    ( 'G2', 2),
                    ( 'G2', 3)],
                    )
```

```
In [38]: 1 df = pd.DataFrame(np.random.randn(6,2),index=hier_index,columns=[ 'A' ,
2 df
```

```
Out[38]:
```

		A	B
G1	1	0.302665	1.693723
	2	-1.706086	-1.159119
	3	-0.134841	0.390528
G2	1	0.166905	0.184502
	2	0.807706	0.072960
	3	0.638787	0.329646

Now let's show how to index this! For index hierarchy we use `df.loc[]`, if this was on the columns axis, you would just use normal bracket notation `df[]`. Calling one level of the index returns the sub-dataframe:

```
In [39]: 1 df.loc[ 'G1' ]
```

```
Out[39]:
```

	A	B
1	0.302665	1.693723
2	-1.706086	-1.159119
3	-0.134841	0.390528

```
In [40]: 1 df.loc['G1'].loc[1]
```

```
Out[40]: A    0.302665  
         B    1.693723  
         Name: 1, dtype: float64
```

```
In [41]: 1 df.index.names
```

```
Out[41]: FrozenList([None, None])
```

```
In [42]: 1 df.index.names = ['Group', 'Num']
```

```
In [43]: 1 df
```

```
Out[43]:
```

		A	B
Group	Num		
G1	1	0.302665	1.693723
	2	-1.706086	-1.159119
	3	-0.134841	0.390528
G2	1	0.166905	0.184502
	2	0.807706	0.072960
	3	0.638787	0.329646

```
In [44]: 1 df.xs('G1')
```

```
Out[44]:
```

	A	B
Num		
1	0.302665	1.693723
2	-1.706086	-1.159119
3	-0.134841	0.390528

```
In [45]: 1 df.xs(['G1', 1])
```

```
Out[45]: A    0.302665  
         B    1.693723  
         Name: (G1, 1), dtype: float64
```

```
In [46]: 1 df.xs(1,level='Num')
```

Out[46]:

	A	B
Group		
G1	0.302665	1.693723
G2	0.166905	0.184502

Great Job!

```
In [ ]: 1
```