Logistic Regression Project

In this project we will be working with a fake advertising data set, indicating whether or not a particular internet user clicked on an Advertisement. We will try to create a model that will predict whether or not they will click on an ad based off the features of that user.

This data set contains the following features:

- 'Daily Time Spent on Site': consumer time on site in minutes
- 'Age': cutomer age in years
- 'Area Income': Avg. Income of geographical area of consumer
- 'Daily Internet Usage': Avg. minutes a day consumer is on the internet
- 'Ad Topic Line': Headline of the advertisement
- 'City': City of consumer
- 'Male': Whether or not consumer was male
- 'Country': Country of consumer
- 'Timestamp': Time at which consumer clicked on Ad or closed window
- 'Clicked on Ad': 0 or 1 indicated clicking on Ad

Import Libraries

Import a few libraries you think you'll need (Or just import them as you go along!)

In [1]:
```
1
```

In [ ]:
```
1
```

Get the Data

Read in the advertising.csv file and set it to a data frame called ad_data.

In [2]:
```
1
```

Check the head of ad_data

In [3]:
```
1
```

Out[3]:

	Daily Time Spent on Site	Age	Area Income	Daily Internet Usage	Ad Topic Line	City	Male	Country	Timestamp	Clicke on A
0	68.95	35	61833.90	256.09	Cloned 5thgeneration orchestration	Wrightburgh	0	Tunisia	2016-03-27 00:53:11	
1	80.23	31	68441.85	193.77	Monitored national standardization	West Jodi	1	Nauru	2016-04-04 01:39:02	
2	69.47	26	59785.94	236.50	Organic bottom-line service-desk	Davidton	0	San Marino	2016-03-13 20:35:42	
3	74.15	29	54806.18	245.89	Triple-buffered reciprocal time-frame	West Terrifurt	1	Italy	2016-01-10 02:31:19	
4	68.37	35	73889.99	225.58	Robust logistical utilization	South Manuel	0	Iceland	2016-06-03 03:36:18	

Use info and describe() on ad_data

In [4]:
```
1
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 10 columns):
Daily Time Spent on Site    1000 non-null float64
Age                         1000 non-null int64
Area Income                 1000 non-null float64
Daily Internet Usage        1000 non-null float64
Ad Topic Line               1000 non-null object
City                        1000 non-null object
Male                        1000 non-null int64
Country                     1000 non-null object
Timestamp                   1000 non-null object
Clicked on Ad               1000 non-null int64
dtypes: float64(3), int64(3), object(4)
memory usage: 78.2+ KB
```

In [5]:
```
1
```

Out[5]:

	Daily Time Spent on Site	Age	Area Income	Daily Internet Usage	Male	Clicked on Ad
count	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000	1000.00000
mean	65.000200	36.009000	55000.000080	180.000100	0.481000	0.50000
std	15.853615	8.785562	13414.634022	43.902339	0.499889	0.50025
min	32.600000	19.000000	13996.500000	104.780000	0.000000	0.00000
25%	51.360000	29.000000	47031.802500	138.830000	0.000000	0.00000
50%	68.215000	35.000000	57012.300000	183.130000	0.000000	0.50000
75%	78.547500	42.000000	65470.635000	218.792500	1.000000	1.00000
max	91.430000	61.000000	79484.800000	269.960000	1.000000	1.00000

Exploratory Data Analysis

Let's use seaborn to explore the data!

Try recreating the plots shown below!

Create a histogram of the Age

In [14]:
```
1
```

Out[14]:
```
(array([ 12.,   6.,  32.,  21.,  27.,  70.,  48.,  87.,  60.,  38.,  8
2.,
         39.,  86.,  37.,  30.,  68.,  26.,  44.,  30.,  18.,  29.,  1
8.,
         32.,  15.,  10.,  16.,   2.,  10.,   2.,   5.]),
 array([ 19. ,  20.4,  21.8,  23.2,  24.6,  26. ,  27.4,  28.8,  30.2,
         31.6,  33. ,  34.4,  35.8,  37.2,  38.6,  40. ,  41.4,  42.8,
         44.2,  45.6,  47. ,  48.4,  49.8,  51.2,  52.6,  54. ,  55.4,
         56.8,  58.2,  59.6,  61. ]),
 <a list of 30 Patch objects>)
```

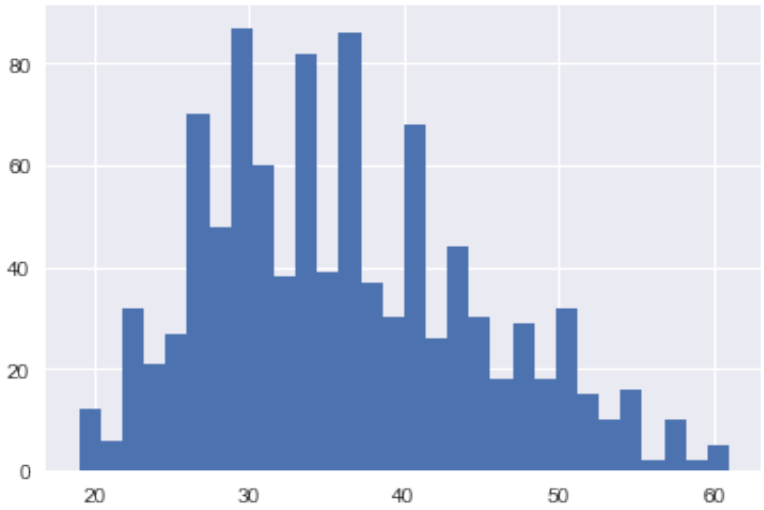

Create a jointplot showing Area Income versus Age.

In [7]:
```
1
```

Out[7]: <seaborn.axisgrid.JointGrid at 0x10ca1e450>

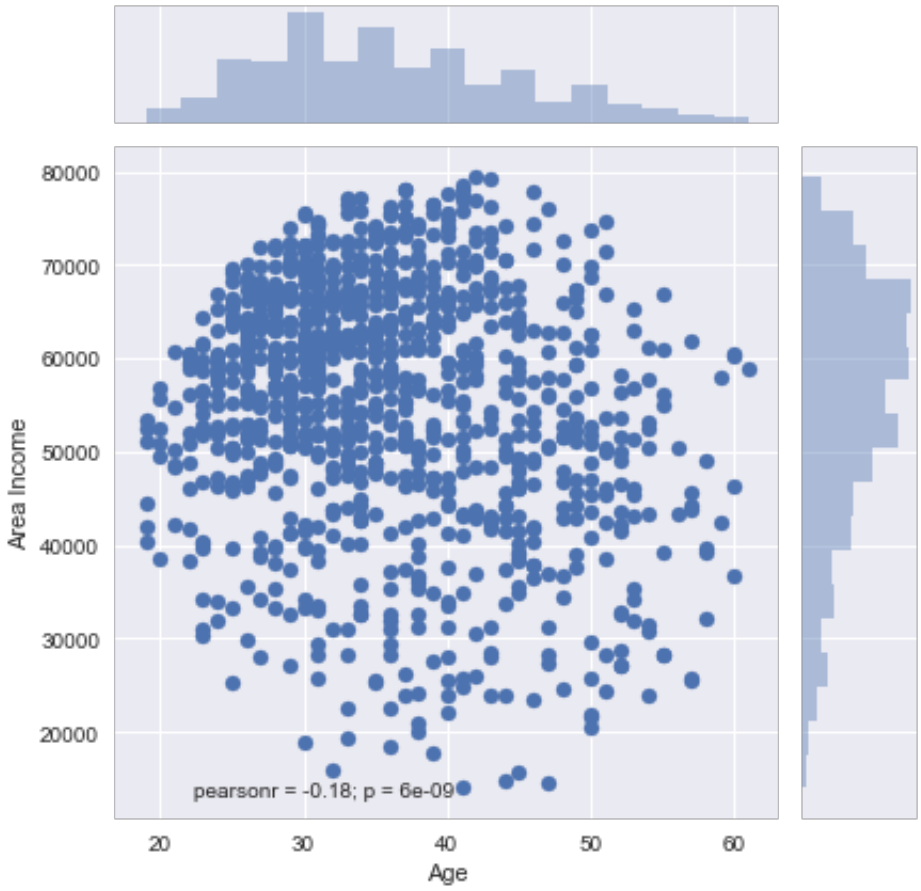

In [ ]:
```
1
```

Create a jointplot showing the kde distributions of Daily Time spent on site vs. Age.

In [13]:
```
1
```

Out[13]: <seaborn.axisgrid.JointGrid at 0x110b75d50>

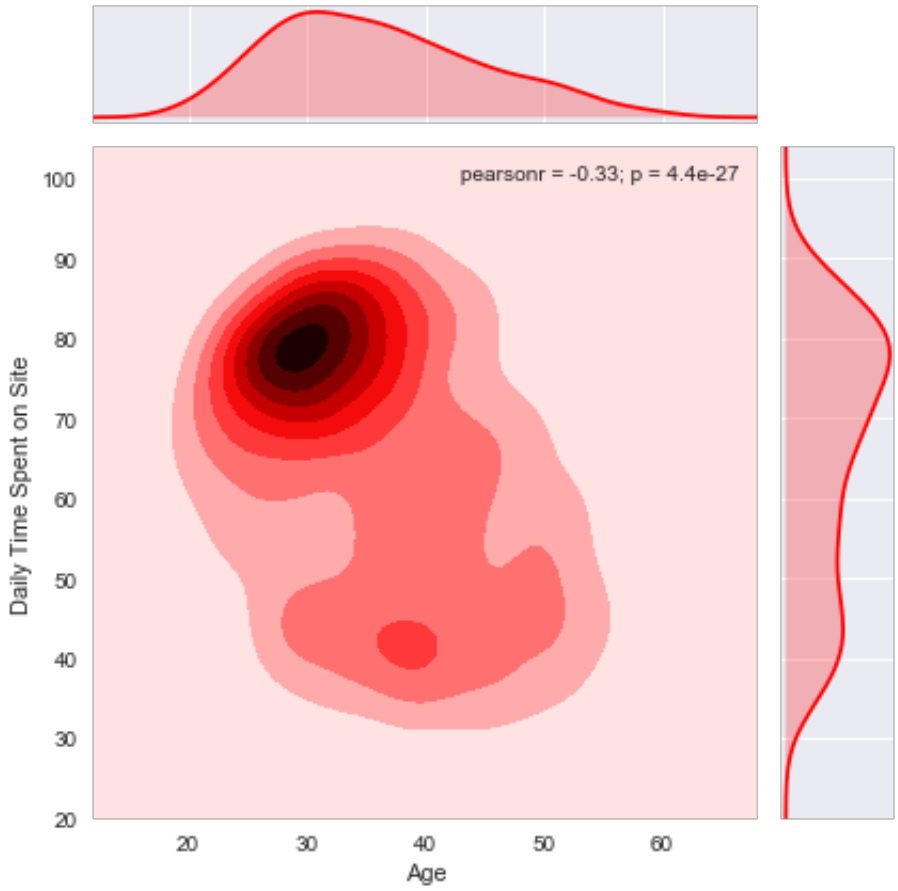

In [ ]:
```
1
```

Create a jointplot of 'Daily Time Spent on Site' vs. 'Daily Internet Usage'

In [15]:
```
1
```

Out[15]: <seaborn.axisgrid.JointGrid at 0x10ffd3f50>

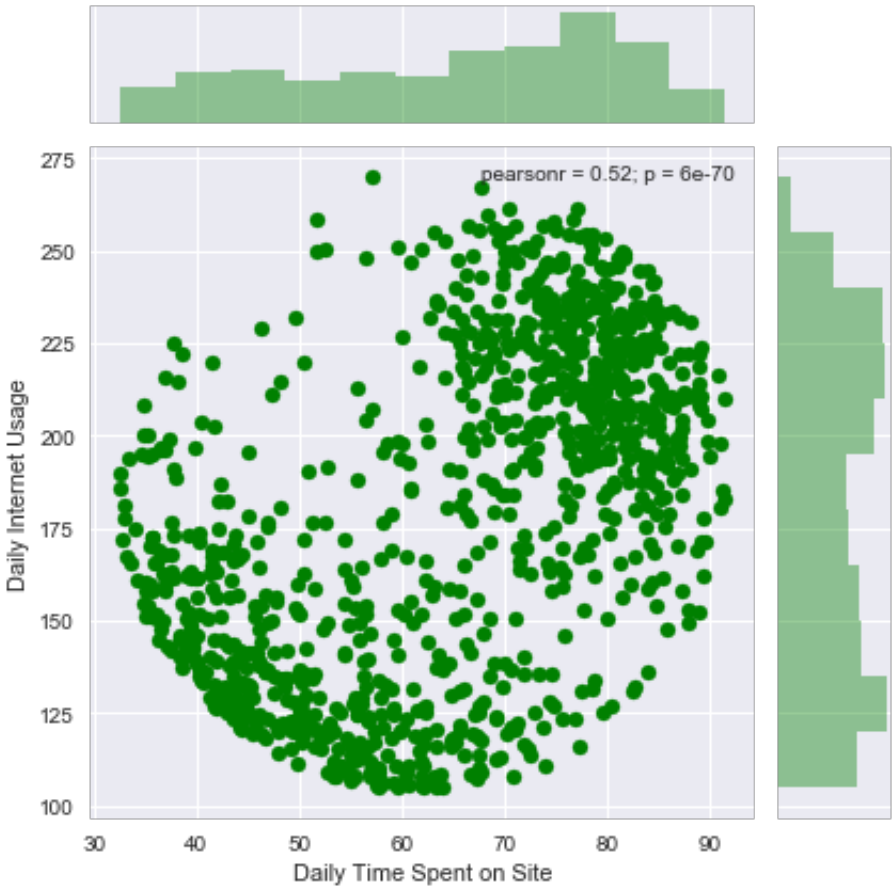

Finally, create a pairplot with the hue defined by the 'Clicked on Ad' column feature.

In [17]:
```
1
```

Out[17]: <seaborn.axisgrid.PairGrid at 0x113787090>

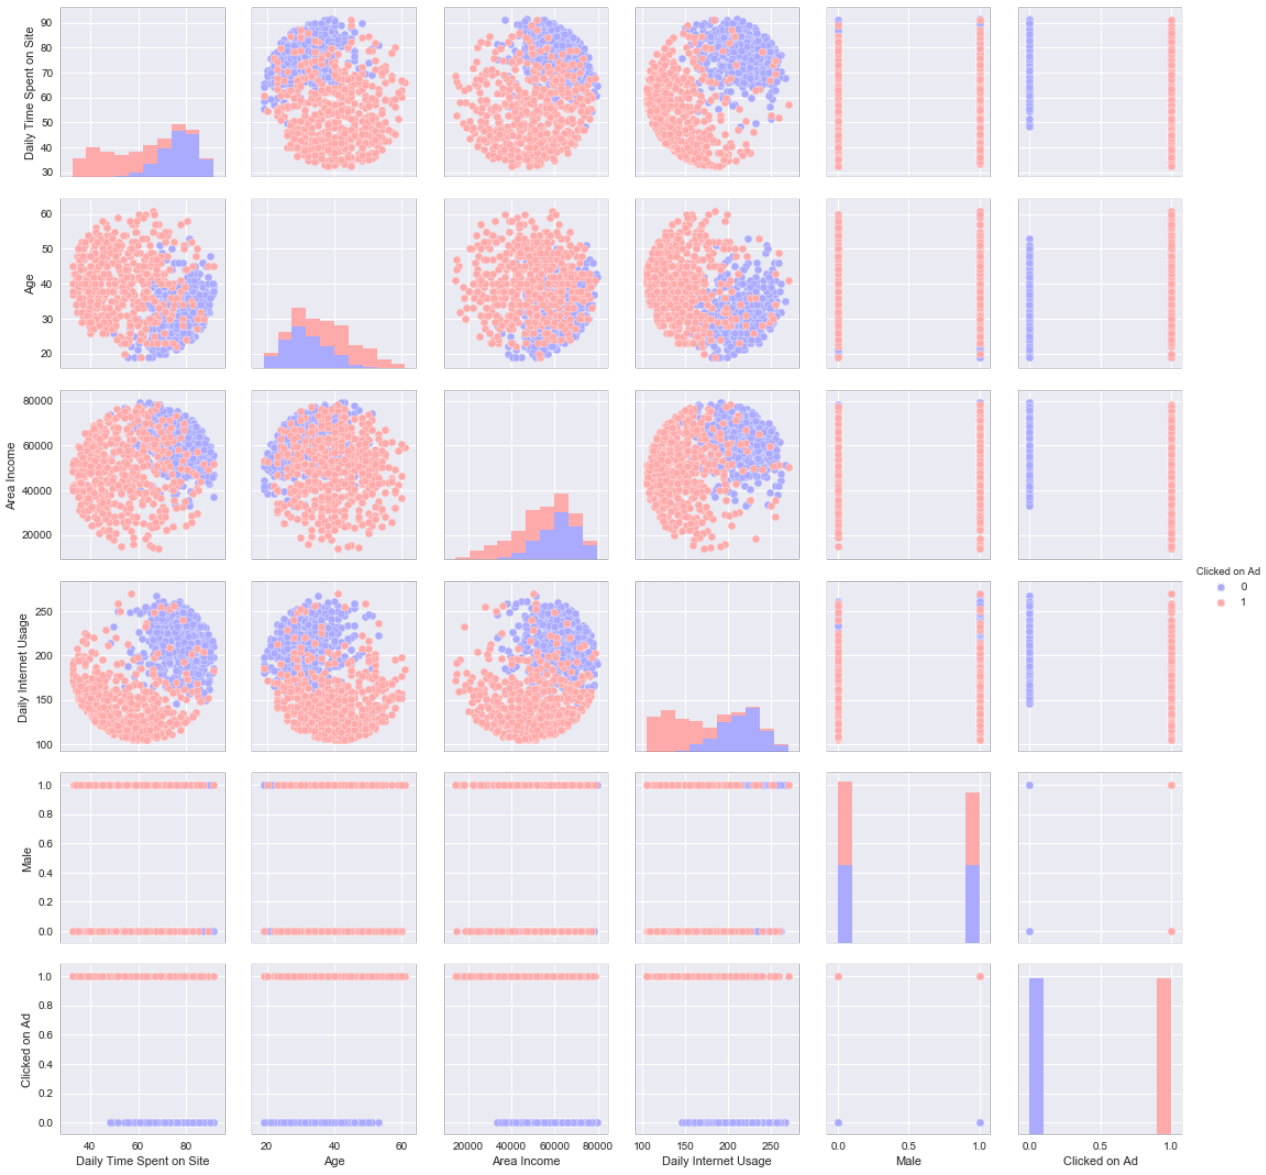

Logistic Regression

Now it's time to do a train test split, and train our model!

You'll have the freedom here to choose columns that you want to train on!

Split the data into training set and testing set using train_test_split

In [38]:
```
1
```

In [39]:
```
1
```

In [47]:
```
1
```

Train and fit a logistic regression model on the training set.

In [48]:
```
1
```

In [49]:
```
1
```

In [50]:
```
1
```

Out[50]:
```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept
=True,
          intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs
=1,
          penalty='l2', random_state=None, solver='liblinear', tol=0.0
001,
          verbose=0, warm_start=False)
```

Predictions and Evaluations

Now predict values for the testing data.

In [51]:
```
1
```

Create a classification report for the model.

In [52]:
```
1
```

In [53]:
```
1
```

```
             precision    recall  f1-score   support

          0       0.87      0.96      0.91       162
          1       0.96      0.86      0.91       168

avg / total       0.91      0.91      0.91       330
```

Great Job!